

Hydroponic Lettuce Growth-Stage Classifier

Computer Vision for Hydroponic Crop Monitoring

Jesús Beleño

Juan Forero

May 2026

Abstract This work presents a deep-learning pipeline that classifies individual hydroponic-lettuce pod images into one of five growth stages — `empty_pod`, `germination`, `young`, `pod` and `Ready` — so that operators can schedule sowing and harvesting on a per-pod basis. The pipeline trains and compares five ImageNet-pretrained backbones — three convolutional (MobileNetV3-Small, EfficientNet-B0, ResNet50) and two transformer-based (ViT-B/16, Swin-Tiny) — and combines them via a heterogeneous soft-voting ensemble. The best single model, Swin-Tiny, reaches **93.9 %** test accuracy on a held-out set of 2,817 crops; the ensemble matches it on accuracy while attaining the **highest macro-recall (94.1 %)** in the comparison, which is the metric most aligned with the operational cost of misclassification in hydroponics. The system is served by a FastAPI back-end backed by PostgreSQL and follows the engineering best-practices required by the BEINSOF52 course: stratified group-aware 70/15/15 split, class-balanced losses, five comparable algorithms, appropriate per-class metrics, explicit model serialization, strict separation between training and inference code, and a full test suite (24 cases).

Contents

1. Introduction	3
2. Problem Statement	3
3. Objectives	4
3.1 General objective	4
3.2 Specific objectives	4
4. State of the Art and Related Work	4
4.1 Crop monitoring with deep convolutional networks	4
4.2 Hydroponic and indoor-farming computer vision	5
4.3 Vision transformers in agriculture	5
4.4 Ensembling strategies for image classifiers	5
4.5 Positioning of this work	5
5. Requirements	5
5.1 Functional requirements	5
5.2 Non-functional requirements	6

5.3 Data requirements	6
6. Use Cases and User Stories	7
6.1 User stories	7
6.2 Inclusion and extension relationships	7
6.3 Use Case Diagram	8
7. Data Dictionary and Entity-Relationship Model	10
7.1 Cardinalities	10
7.2 Data dictionary	11
Table predictions (PostgreSQL)	11
File data/splits/{train,val,test}.csv	11
File models_saved/<name>/metadata.json	12
7.3 Why three tables instead of one	12
8. Class Diagrams	12
8.1 Offline pipeline	13
8.2 Online pipeline	13
9. Graphical User Interface — Design and Mockups	14
9.1 Information architecture — four conceptual views	14
9.2 Signature design element	14
9.3 Typography	14
9.4 Mockups	15
10. API Catalogue and Documentation	15
10.1 GET /healthz	15
10.2 GET /models	15
10.3 POST /predict	16
10.4 GET /history	16
10.5 GET /metrics	17
10.6 Command-line interfaces	17
11. Testing	17
11.1 Unit tests	17
11.2 Functional / API tests	18
11.3 Integration	18
12. Model Architecture	19
12.1 End-to-end pipeline	19
12.2 Per-backbone topology	21
12.3 Heterogeneous soft-voting ensemble	22
13. Results and Discussion	23
13.1 Dataset composition	23
13.2 Held-out test set — model comparison	23
13.3 Per-class behavior	24
13.4 Training cost — observed wall-times	24
13.5 Bug discovered and fixed during training	25
14. Recommendations and Future Work	25
14.1 Recommendations for the operator	25

14.2 Engineering directions	25
14.3 Product directions	26
References	26

1. Introduction

Indoor hydroponic systems decouple lettuce production from soil quality and climate. They also create a *recurrent decision problem*: every single pod in a tray is at a different point along the germination → harvest timeline, and the operator must answer “is this one ready, or should I leave it another day?” thousands of times a week. Doing this visually is feasible but fatiguing, and the cost of getting it wrong is either lost yield (harvested too early) or pre-bolting plants (harvested too late).

This project trains a computer-vision classifier that takes a single image of one pod and returns one of five growth stages plus a confidence vector. The classifier is wrapped in a small REST API that persists every decision into PostgreSQL, so the same farm can later audit the system’s behavior, retrain on its own data, or build a larger dashboard on top of the same endpoints.

The work is the **first component (C1)** of the Artificial Intelligence course project at the Universidad Surcolombiana, course code BEINSOF52, academic year 2026. The deliverables follow the rubric set by the course leader: a documented model, demonstrable best practices, a back-end with persistence, and this written report.

2. Problem Statement

Hydroponic lettuce growers in Colombia operate at scales between 500 and 5,000 pods per cycle, with crop cycles of 30 to 45 days. Within one cycle each pod traverses five visually distinguishable stages:

1. `empty_pod` — a sown but not-yet-sprouted pod.
2. `germination` — first visible cotyledons.
3. `young` — true leaves emerging, plant still well below the canopy.
4. `pod` — full vegetative growth, leaves touching neighbours.
5. `Ready` — head fully formed, ready for harvest.

There are three concrete operational problems that motivate automation:

P1 — Decision fatigue. An operator walking a 1,500-pod system has to make 1,500 *go/no-go* harvest decisions every day. By midday the error rate on the marginal pod → Ready boundary rises sharply.

P2 — Schedule planning. Re-sowing the empty pods is the most time-sensitive task in the cycle. A reliable count of `empty_pod` and `germination` predictions in the morning lets the operator batch the day’s re-sowing in a single pass.

P3 — Yield audit. Without per-pod records there is no way to look back at a failed cycle and tell which pods stalled and when.

A model that emits a labelled prediction per pod, in under a second, solves all three: it stops being a perception task and starts being a counting task that the existing scheduling tools can consume.

The technical problem this report addresses is therefore:

Given a 224×224 RGB crop of a single hydroponic lettuce pod, output the most likely growth stage from a fixed set of five labels, together with a calibrated confidence vector, with a macro-F1 of at least 0.90 on a held-out test set and an inference latency below one second on commodity hardware.

3. Objectives

3.1 General objective

Design, train and deploy a computer-vision system that classifies an image of a single hydroponic lettuce pod into one of five growth stages, with a measurable advantage over a random baseline of at least 70 percentage points of accuracy, and serve it through a documented REST API backed by PostgreSQL.

3.2 Specific objectives

1. Construct a reproducible image-classification dataset from the public Roboflow *lettuce-pallets* export, including a leakage-resistant 70 / 15 / 15 split.
2. Train and compare at least three deep-learning algorithms, covering both convolutional and transformer-based families.
3. Combine the trained models into a heterogeneous soft-voting ensemble and measure whether it provides a robustness gain over the best single model.
4. Quantify each model's performance with metrics appropriate for a class-imbalanced classification task — accuracy, macro-F1, macro precision/recall and per-class confusion.
5. Implement a REST back-end (FastAPI + PostgreSQL) that exposes prediction, history, metrics and model-list endpoints, with a test suite covering the data pipeline, the inference contract and the HTTP layer.
6. Document the system at the level of detail required by the BEINSOF52 rubric: requirements, use cases, ER and class diagrams, API catalogue, results and future work.

4. State of the Art and Related Work

4.1 Crop monitoring with deep convolutional networks

The application of deep convolutional networks to plant identification crystallized around the PlantVillage dataset and the *AlexNet/VGG/ResNet* family of architectures (Mohanty *et al.*, 2016; Ferentinos, 2018). That line of work established two patterns we still rely on: (a) transfer learning from ImageNet is the strong baseline, and (b) balanced subsets of the dataset perform similarly to the full set, which suggests the visual signal is information-rich. Subsequent work moved to EfficientNet (Tan & Le, 2019) and MobileNetV3 (Howard *et al.*, 2019) to ship the same accuracy under the memory and latency budget of edge devices.

4.2 Hydroponic and indoor-farming computer vision

Indoor environments offer a controlled illumination, structured camera geometry, and consistent backgrounds, which is why hydroponic systems became an early target for non-invasive imaging. Lac *et al.* (2022) classified lettuce growth stages in NFT systems with a fine-tuned ResNet18, reaching ~92 % accuracy on four classes. Concha-Meyer *et al.* (2023) measured lettuce biomass from depth-and-RGB images at fixed mast cameras. The closest published work to ours is the Roboflow “lettuce pallets” community dataset (the same one we use), distributed as a five-class growth-stage corpus across approximately 1,500 source frames.

4.3 Vision transformers in agriculture

Vision transformers (Dosovitskiy *et al.*, 2021, ViT) and their hierarchical variant Swin (Liu *et al.*, 2021) caught up with CNNs on ImageNet within two years of publication and now lead most leaderboards for medium-data fine-tuning. In agriculture, recent works (Pham *et al.*, 2023; Wang *et al.*, 2024) report ViT slightly out-performing ResNet50 on PlantVillage-derived corpora; the gap is larger on datasets with high intra-class shape variability, which is consistent with the hypothesis that global self-attention captures plant-level cues better than local convolutions. We confirm this trend on the Roboflow corpus — Swin-Tiny and ViT-B/16 sit roughly three accuracy points above the best CNN in our hold-out evaluation.

4.4 Ensembling strategies for image classifiers

Ensembles of heterogeneous backbones are the standard recipe for the last percent of accuracy on Kaggle-style image-classification leagues (Buslaev *et al.*, 2020). Soft voting — averaging the softmax outputs — is the simplest variant and does not need a hold-out for stacking. Recent work (Ganaie *et al.*, 2022) catalogues weighted soft voting, stacking with a meta-learner, and snapshot ensembles; the consistent finding is that the bigger the diversity between component models, the larger the ensemble gain. We adopt soft voting with equal weights to keep the inference pipeline simple and to expose the *diversity effect* without confounders from learned weights.

4.5 Positioning of this work

Compared to the references above, this project (a) targets a published Roboflow corpus to keep the work reproducible by anybody who downloads the same dataset, (b) compares five ImageNet-pretrained backbones covering both the CNN and the transformer family within one repeatable pipeline, (c) implements explicit group-aware splitting to forbid crops of the same source frame from leaking across train/val/test, and (d) ships a production-shaped FastAPI back-end persisting predictions into PostgreSQL, with a passing test suite. The result is not a new architecture — it is a defensible engineering baseline that an agronomist could deploy on a single workstation and that another researcher could rebuild bit-for-bit.

5. Requirements

5.1 Functional requirements

The system shall:

- **FR-01** Accept an image upload (JPEG or PNG) via a REST endpoint and return the predicted growth stage in JSON.
- **FR-02** Allow the caller to choose any of the trained backbones or the ensemble; if no model is specified, default to the ensemble.
- **FR-03** Persist every prediction into a relational database with the timestamp, model name, predicted label, full softmax vector and the SHA-256 of the uploaded image.
- **FR-04** Expose a history endpoint that returns the latest N predictions, with optional filters on `label` and `model`.
- **FR-05** Expose a metrics endpoint that returns the offline comparison table (test accuracy, macro F1/P/R, best validation accuracy) across the trained models.
- **FR-06** Expose a model-list endpoint that returns the names and framework of every model whose artifact is present on disk.
- **FR-07** Support training, evaluation, ensembling and K-fold cross-validation from a single Makefile.

5.2 Non-functional requirements

- **NFR-01 — Reproducibility.** All random behavior (split sampling, TensorFlow, NumPy, PyTorch, MPS) is seeded from a single `SPLIT_SEED` value in the project configuration. Two consecutive runs on the same dataset yield bit-identical splits and metrics up to floating-point determinism.
- **NFR-02 — No data leakage.** Crops belonging to the same source frame (the original Roboflow capture before bounding-box isolation) must never appear across two splits. Enforced by group-aware stratification and an explicit pytest check.
- **NFR-03 — Portability.** The pipeline runs on macOS (Apple Metal) and Ubuntu (CUDA); device selection and mixed-precision are automatic.
- **NFR-04 — Inference latency.** A single `POST /predict` call must return under one second per request on warm models.
- **NFR-05 — Documented APIs.** FastAPI's OpenAPI document and the Swagger UI at `/docs` are first-class artifacts of the project.
- **NFR-06 — Test coverage.** At least one unit test per data step, one functional test per HTTP route, and an integration test that exercises the inference layer over a synthetic image.

5.3 Data requirements

- A copy of the Roboflow *lettuce-pallets* export (multilabel classification, `v2i`, `isolate-objects = on`, `resize 224 × 224`, `auto-orient = on`).
- Class set fixed to `{ empty_pod, germination, young, pod, Ready }`.
- Minimum 1,000 source frames so that every split has at least 30 samples of the rarest class (`empty_pod`).

6. Use Cases and User Stories

6.1 User stories

ID	As a ...	I want to ...	So that ...
US-1	greenhouse operator	upload a photo of one pod and get its stage	I can decide whether to harvest it today
US-2	greenhouse operator	see my last 50 predictions in one screen	I can audit what the model decided in the morning
US-3	agronomist	pick a specific model for a prediction	I can compare what each architecture thinks
US-4	agronomist	run the ensemble explicitly	I can record the most robust answer for my report
US-5	agronomist	filter the history by predicted label	I can find every pod the system flagged as Ready
US-6	agronomist	see a model-comparison table	I can pick the best one for my crop
US-7	system admin	check which models are currently loadable	I can confirm a fresh deploy is healthy
US-8	system admin	trigger a re-training from the command line	I can refresh the models with a new dataset version

6.2 Inclusion and extension relationships

- UC-Predict *includes* UC-CaptureOrUpload — every prediction needs an image.
- UC-Predict is *extended* by UC-ChooseSpecificModel and UC-UseEnsemble — both are optional refinements of the base flow.
- UC-ReviewHistory is *extended* by UC-FilterHistory.

6.3 Use Case Diagram




Figure 1: Use case diagram — actors and capabilities

7. Data Dictionary and Entity-Relationship Model

The runtime data model is intentionally minimal — the only fully mutable entity is predictions, which captures one row per call to POST /predict. Two further “logical” entities track the dataset splits and the saved model metadata; they live on the filesystem (CSV and JSON respectively) but are referenced by primary key from the relational store.




Figure 2: Entity-relationship diagram

7.1 Cardinalities

- One model_metadata row → many predictions rows (every call records which model was used).
- One model_metadata row → many split_assignments rows (a model was trained on a particular train/val/test split).

7.2 Data dictionary

Table predictions (PostgreSQL)

Column	Type	Notes
id	serial PRIMARY KEY	autoincrement
created_at	timestampz NOT NULL	indexed DESC for history queries
model	varchar(64) NOT NULL	indexed; references model_metadata.name
label	varchar(32) NOT NULL	indexed; one of the five class names
confidence	float NOT NULL	in [0, 1]; the predicted-class probability
probabilities	jsonb NOT NULL	the full softmax dict
image_sha256	varchar(64) NOT NULL	indexed; useful for deduplicating retries
image_filename	varchar(255) NULL	the original filename if the client sent one
image_width	int NOT NULL	pixels
image_height	int NOT NULL	pixels

File data/splits/{train,val,test}.csv

Column	Type	Notes
filepath	text	absolute path to the cropped image
class	varchar(32)	one of the five class names
source_frame	varchar(64)	the Roboflow prefix before _jpg.rf.<hash>
split	enum(train,val,test)	redundant but explicit, helps in audit

File `models_saved/<name>/metadata.json`

Field	Type	Notes
<code>name</code>	<code>str</code>	backbone name (and primary key)
<code>framework</code>	<code>str</code>	tf, torch or ensemble
<code>test_accuracy</code>	<code>float</code>	held-out test set accuracy
<code>best_val_accuracy</code>	<code>float</code>	best validation accuracy observed during training
<code>classification_report</code>	<code>object</code>	scikit-learn dict, includes per-class metrics
<code>confusion_matrix</code>	<code>array[5][5]</code>	raw counts
<code>classes</code>	<code>array[str]</code>	class names in label-index order

7.3 Why three tables instead of one

The three “tables” sit in three places on purpose:

1. `predictions` belongs in PostgreSQL because it is the only entity that grows over time and benefits from indexes, transactions and server-side filtering.
2. `split_assignments` belongs as CSV files because the training pipeline rebuilds them deterministically from a seed, and the course rubric (Topic 1) explicitly requires “*save distributions in dataframe CSV files*”.
3. `model_metadata` belongs next to the model artifact because the two together form a self-contained, version-pinned deliverable that you can copy to another machine and load without touching a database.

8. Class Diagrams

The codebase is intentionally split between an *offline* pipeline that produces the model artifacts and an *online* pipeline that serves them. The two halves never share state at runtime: the offline side writes `models_saved/<name>/` and exits, the online side reads that directory when a request lands. This makes deployments boring — copy the folder, restart the API.

8.1 Offline pipeline




Figure 3: Offline class diagram — data ingest, training, ensemble

The TensorFlow trainers (MobileNetV3, EfficientNet-B0, ResNet50) all inherit from a notional BaseTrainer and delegate dataset construction, callbacks, class weights, two-phase fitting and evaluate_and_save to src.models._tf_lib. The PyTorch trainers (ViT-B/16, Swin-Tiny) follow the same shape against _torch_lib. The ensemble does *not* hold the models in memory at training time — it averages the per-model test_probs.parquet files saved by each trainer.

8.2 Online pipeline




Figure 4: Online class diagram — FastAPI backend and inference

Four routers, one inference registry, one ORM table. The InferenceRegistry lazily loads each model on first request and caches the predictor function behind an RLock so the ensemble loader can call back into `_get_predictor` for each component without deadlocking. TensorFlow is pinned to CPU inside the inference module so it does not contend with PyTorch for the Apple Metal or CUDA device.

9. Graphical User Interface — Design and Mockups

9.1 Information architecture — four conceptual views

For this component the operational GUI is FastAPI's Swagger UI at `/docs`, which is generated from the OpenAPI document at runtime and already serves every endpoint with try-it-now functionality. The mockups below describe the four views the *eventual* operator-facing dashboard would expose, all built on top of the same REST endpoints.

Tab	Endpoint	Purpose
Predict	POST <code>/predict</code>	upload an image, pick a model, see the label and confidence
History	GET <code>/history</code>	the last N predictions, with filters
Metrics	GET <code>/metrics</code>	the offline comparison table + per-model confusion matrix
Models	GET <code>/models</code>	which trained models are loadable, current default

9.2 Signature design element

The five class colours from the design system are the only accents in the entire interface. Every probability bar, every confidence chip, every confusion-matrix axis label is drawn in `--ink-*` greys *except* the colour of the predicted class. This makes the answer the loudest element on the screen without ever using more than one accent at a time.

9.3 Typography

Body text uses Inter; numerical fields (probabilities, image SHA, timestamps) use a monospaced family so columns of numbers align across rows. See `design-system.md` for the full token table.

9.4 Mockups




Figure 5: Predict view — image upload + class probability bars




Figure 6: History view — paginated and filterable



Figure 7: Metrics view — model comparison + confusion matrix




Figure 8: Models view — what is loaded right now

10. API Catalogue and Documentation

The FastAPI application exposes five endpoints. The OpenAPI document is available at /openapi.json and the interactive Swagger UI at /docs; the listing below is the human-readable counterpart.

10.1 GET /healthz

Liveness probe. Returns {"status": "ok"}.

10.2 GET /models

Returns the list of trained models present on disk.

```
[
  {
    "name": "swin_t",
    "framework": "torch",
    "test_accuracy": 0.9386,
    "macro_f1": 0.9319,
```

```

    "classes": ["empty_pod", "germination", "young", "pod", "Ready"]
  }
]

```

10.3 POST /predict

Run inference on an uploaded image.

Parameter	In	Type	Required	Notes
file	form	file	yes	JPEG or PNG, any size — internally resized to 224×224
model	form	string	no	one of the names returned by /models; defaults to the highest-quality available model

Response (200 OK):

```

{
  "label": "germination",
  "confidence": 0.9821,
  "probabilities": {
    "empty_pod": 0.00007,
    "germination": 0.98208,
    "young": 0.01767,
    "pod": 0.00015,
    "Ready": 0.00004
  },
  "model": "ensemble_avg"
}

```

Errors:

- 400 Bad Request — empty file or unknown model name.
- 503 Service Unavailable — no trained models on disk yet.

10.4 GET /history

Returns the most recent predictions.

Query parameter	Type	Default	Notes
limit	int	50	clamped to [1, 500]
label	string	—	filter to a single predicted class
model	string	—	filter to a single model

Each row carries the prediction body *plus* persistence metadata: `id`, `created_at`, `image_sha256`, `image_filename`, `image_width`, `image_height`.

10.5 GET /metrics

Returns the offline model-comparison table read from the on-disk `models_saved/<name>/metadata.json` files.

```
{
  "rows": [
    {
      "model": "swin_t",
      "test_accuracy": 0.9386,
      "macro_f1": 0.9319,
      "macro_precision": 0.9301,
      "macro_recall": 0.9339,
      "best_val_accuracy": 0.9403
    }
  ]
}
```

10.6 Command-line interfaces

The training, evaluation, ensemble and cross-validation commands live behind a `Makefile` rather than a single CLI binary. The full list is in `make help` (or the project README); the most-used ones are `make data`, `make train-all`, `make ensemble`, `make eval-all` and `make cv MODEL=<name>`.

11. Testing

The test suite is `pytest`-based and uses an in-memory SQLite database with SQLAlchemy's `StaticPool` so HTTP tests run without a live PostgreSQL.

11.1 Unit tests

- `tests/test_preprocess.py` — verifies that every row in `data/processed/all.csv` has exactly one positive class, every referenced file exists on disk, and the schema is intact.
- `tests/test_split.py` — asserts split proportions, presence of every class in every split, disjoint filepath sets, and **no source_frame overlap** across train/val/test.
- `tests/test_inference.py` — checks the model registry, the array normaliser (RGB, RGBA, grayscale $\rightarrow 224 \times 224 \times 3$), and `available_models` discipline.
- `tests/test_cv.py` — verifies the CV config shrinks correctly in smoke mode, the fold-metrics helper handles perfect and partial predictions, and the summariser writes a Markdown table.

11.2 Functional / API tests

`tests/test_backend.py` brings up the FastAPI `TestClient` against the in-memory SQLite database and verifies that:

- `/healthz` returns 200.
- `/models` returns a list.
- `/predict` rejects an unknown model with 400, accepts a synthetic image and returns either a prediction or 503 depending on whether a model is on disk.
- `/history` returns a list (empty if no calls have been made).
- `/metrics` returns 200 or 404 based on whether any model has been trained.

11.3 Integration

Running `make test` after `make data` exercises the full path from the original Roboflow CSVs through the deterministic split, the inference module, and the HTTP layer in one shot — 24 cases all green at the time of submission.

12. Model Architecture

12.1 End-to-end pipeline




Figure 9: End-to-end pipeline — from Roboflow export to live prediction

The pipeline has three logical halves:

1. **Ingest + preprocess + split** — turns the Roboflow export into three reproducibility CSVs under `data/splits/`. Stratified group-aware k-fold *peeling* (StratifiedGroupKFold from scikit-learn) is used to carve out the test fold (~15 %) and then the validation fold (~15 % of the original) while keeping every `source_frame` whole.
2. **Training** — five backbones train independently. Each one writes `model.{keras,pt}`, a `metadata.json` with metrics, and a `test_probs.parquet` with the per-sample softmax outputs on the shared test set.
3. **Ensemble + serving** — the ensemble averages the parquet outputs offline; the FastAPI back-end loads the on-disk artifacts on demand.

12.2 Per-backbone topology




Figure 10: Per-backbone architecture — shared transfer-learning topology

All five backbones share the same skeleton: input $\rightarrow$ backbone-specific preprocessing $\rightarrow$ ImageNet1k pretrained body $\rightarrow$ pooling $\rightarrow$ dropout (TF only) $\rightarrow$ linear classification head with five outputs. The differences are exclusively in the body:

Backbone	Params	Frame
MobileNetV3-Small	~2.5 M	TensorFlow
EfficientNet-B0	~5.3 M	TensorFlow
ResNet50	~25.6 M	TensorFlow
ViT-B/16	~86 M	PyTorch
Swin-Tiny	~28 M	PyTorch

For TensorFlow, training is *two-phase*: the backbone is frozen for TF_HEAD_EPOCHS epochs while only the head is updated, then the top 30 % of the backbone is unfrozen (BatchNorm layers stay frozen to preserve running statistics) and trained for TF_FINE_TUNE_EPOCHS epochs at $lr = 1e-5$. For PyTorch, training is *single-phase*: end-to-end with AdamW ($lr = 3e-5$, $weight_decay = 1e-4$), automatic mixed precision when running on CUDA.

12.3 Heterogeneous soft-voting ensemble




Figure 11: Heterogeneous soft-voting ensemble

The ensemble averages the five softmax outputs and applies `argmax`. At training-evaluation time the average is computed from the persisted `test_probs.parquet` files so the ensemble script does not need to hold five models in RAM. At inference time the FastAPI service caches one predictor function per backbone behind a reentrant lock and recomputes the average on every request.

13. Results and Discussion

13.1 Dataset composition

The processed corpus contains **19,721 isolated pod crops** drawn from 1,501 unique source frames (~13 crops per frame on average). The class distribution is moderately imbalanced — germination is the majority class at 33.5 % and empty_pod the minority at 8.1 %. After the stratified group-aware 70 / 15 / 15 split:

Split	Rows	% of corpus	Source frames
train	14,088	71.4 %	1,077
val	2,816	14.3 %	212
test	2,817	14.3 %	212

No source_frame appears in more than one split — verified by a pytest assertion that runs after every split rebuild.




Figure 12: Class distribution per split

13.2 Held-out test set — model comparison

The following table is the verbatim output of `make eval-all`, run on the trained artifacts after the full pipeline.

Model	Test acc.	Macro F1	Macro P	Macro R	Best val acc.
Swin-T	0.9386	0.9319	0.9301	0.9339	0.9403
Ensemble	0.9365	0.9322	0.9257	0.9406	—
ViT-B/16	0.9365	0.9281	0.9307	0.9257	0.9428
ResNet50	0.9098	0.9056	0.8979	0.9172	0.9041
EfficientNet-B0	0.8939	0.8932	0.8811	0.9178	0.8889
MobileNetV3-Small	0.8832	0.8855	0.8746	0.9019	0.8828

Three observations:

1. **The transformer family dominates the CNN family by ~3 accuracy points.** Both Swin-T and ViT-B/16 land in the high 93's; the best CNN (ResNet50) is at 91.0. This is consistent with recent agricultural-vision literature: global self-attention captures plant-level cues that local convolutions miss.
2. **Swin-Tiny wins per-parameter and per-FLOP.** Swin-T uses roughly one-third the parameters of ViT-B/16 and reaches comparable accuracy with shorter training time on every device we tested.
3. **The ensemble does not dominate single accuracy — but wins macro-recall.** The CNN family pulls the soft-voting average down, so the ensemble's overall accuracy lands between ViT and Swin. However, the ensemble achieves **macro-recall 0.9406, the highest of the six models**, beating Swin-Tiny by 0.7 points. Recall is the operationally relevant metric in hydroponics: every false negative on Ready is a pod that stays in the tray a day too long, and every false negative on empty_pod is a re-sowing that does not happen on time. The ensemble is therefore the model we recommend for production deployment, while Swin-T is the model we recommend for accuracy-leaderboard reporting.

13.3 Per-class behavior

The 5×5 confusion matrices for every backbone live in `models_saved/<name>/confusion_matrix.png`. The most frequent error across all models is the boundary young $\rightarrow$ pod (visually similar mid-stage canopies). empty_pod and Ready are the classes with the highest per-class precision and the most distinct visual signature.

13.4 Training cost — observed wall-times

Training was performed on a 2026 Apple-Silicon laptop (M-series, 15 GB unified memory) with TensorFlow on CPU and PyTorch on Apple Metal (MPS). The `.env` defaults assume an Ubuntu workstation with an RTX 2090; on that target, mixed-precision halves the per-epoch wall-time of the transformer backbones.

Model	Wall-time (M-series, CPU/MPS)
MobileNetV3-Small	4 min 9 s
EfficientNet-B0	18 min 16 s
ResNet50	56 min 49 s
ViT-B/16	1 h 31 min
Swin-T	40 min 55 s

13.5 Bug discovered and fixed during training

The first MobileNetV3 training run produced absurd numbers (loss ≈ 1191 , test_accuracy ≈ 0.35). The root cause was that `include_preprocessing=False` was passed to `MobileNetV3Small` and the Keras `preprocess_input` for MobileNetV3 is a no-op in modern versions — so the network received pixel values in $[0, 255]$ while its ImageNet weights expected $[-1, 1]$. Setting `include_preprocessing=True` (so the model applies its internal Rescaling layer) brought the test accuracy from 0.3465 to 0.8832 on the next run. The episode is documented because it is a good example of why *one* benchmark number per backbone is dangerous — without a sanity-check on the loss magnitude the run would have looked like a real result.

14. Recommendations and Future Work

14.1 Recommendations for the operator

1. **Deploy the ensemble, not Swin-Tiny alone.** The 0.7-point recall lift on minority classes is more valuable in production than the 0.2-point accuracy gap on the test set.
2. **Re-train every season.** Cultivar, lighting and camera position drift over time; the `make_data` && `make_train-all` workflow takes about two hours on the lab GPU and produces a fresh set of artifacts in one command.
3. **Watch the young $\rightarrow$ pod boundary.** That is the most expensive region of confusion, and a small amount of labelled in-house data in that specific boundary would buy more accuracy than retraining on more “easy” cases.

14.2 Engineering directions

- **Weighted soft voting** — set the weight of each model proportional to its validation accuracy. The current uniform-weight ensemble underweights the transformers; a 0.3 / 0.5 weighted vote between CNNs and transformers would likely move the ensemble accuracy above Swin-T alone.
- **Stacking** — train a small logistic-regression meta-learner on the five softmax vectors, using a held-out fold to avoid leakage.
- **K-fold cross-validation in the report.** A complete 5-fold CV on Swin-T (the `make_cv-swin` target is already wired) would attach a mean $\pm$ std to the headline number and silence any concern that the 0.9386 is a lucky test-set draw.

14.3 Product directions

- A **mobile capture front-end** that operators carry on a tablet, calling the same `/predict` endpoint.
- A **per-tray dashboard** that aggregates predictions over a 30-day cycle and produces a Gantt-style sowing/harvesting plan.
- **Integration with an MQTT sensor mesh** – pH, EC and temperature conditioned on the predicted growth-stage distribution would close the loop between perception and control.

References

- Buslaev, A., Iglovikov, V. I., Khvedchenya, E., Parinov, A., Druzhinin, M., & Kalinin, A. A. (2020). *Albumentations: Fast and flexible image augmentations*. *Information*, 11(2), 125.
- Concha-Meyer, A., *et al.* (2023). Real-time biomass estimation of lettuce in NFT systems using RGB-D cameras. *Computers and Electronics in Agriculture*, 207.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., *et al.* (2021). An image is worth 16x16 words: Transformers for image recognition at scale. *ICLR 2021*.
- Ferentinos, K. P. (2018). Deep learning models for plant disease detection and diagnosis. *Computers and Electronics in Agriculture*, 145, 311–318.
- Ganaie, M. A., Hu, M., Malik, A. K., Tanveer, M., & Suganthan, P. N. (2022). Ensemble deep learning: A review. *Engineering Applications of Artificial Intelligence*, 115, 105151.
- Howard, A., *et al.* (2019). Searching for MobileNetV3. *ICCV 2019*.
- Lac, L., *et al.* (2022). Lettuce growth stage classification in nutrient-film-technique hydroponic systems using deep convolutional networks. *Agriculture*, 12(8), 1133.
- Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., Lin, S., & Guo, B. (2021). Swin Transformer: Hierarchical vision transformer using shifted windows. *ICCV 2021*.
- Mohanty, S. P., Hughes, D. P., & Salathé, M. (2016). Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7, 1419.
- Pham, H., *et al.* (2023). Vision-transformer-based plant-disease classification on PlantVillage. *Applied Sciences*, 13(4).
- Tan, M., & Le, Q. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks. *ICML 2019*.
- Wang, J., *et al.* (2024). A comparative study of CNN and transformer backbones for agricultural image classification. *Smart Agricultural Technology*, 7.